

DEEP-TECH MARKET RESEARCH

Method and Limitations

Physical AI in Switzerland and Europe

Lucas Bergerot

18 August 2026

Deliverable

This document is the actual deliverable of the project. Everything downstream is execution. Nothing gets coded until every section below is filled in and defensible out loud.

1. Sub-sector definition

Sector: Physical AI. Companies whose core product is a machine that perceives and acts in the physical world, and whose value comes primarily from onboard autonomy rather than from the mechanics. Mobile robots, drones, autonomous industrial machines, manipulation, humanoids.

Geography: Switzerland and Europe (EU + UK + Norway), by headquarters.

In scope — a company qualifies if all four hold: - Its main commercial product is a physical system, and the differentiator is the AI on board - It is independently financed — a startup or scale-up, not a division of a listed group - Headquarters in Switzerland or Europe - At least one disclosed funding round, or public evidence of revenue

Out of scope — explicitly excluded, with the reason: - **Pure software, simulation or foundation-model companies** — no physical product, so the same criteria would not measure the same thing - **Industrial automation incumbents** (ABB, KUKA, Stäubli) — mature and orders of magnitude larger; including them compresses every startup into the bottom of the scale - **Swiss and European R&D sites of foreign tech groups** (Google Zurich, Apple, NVIDIA, Meta, Disney Research) — scoring Alphabet on "amount raised" is meaningless. Listed unscored in the annex, section 7 - **Academic labs and unincorporated spin-off projects** — no company to score yet

Edge case rulings, decided in advance: - *Founded in Europe, HQ later moved to the US* → out of the main ranking, listed in the annex and flagged. The migration is itself a finding worth discussing in the thesis. - *European HQ, manufacturing in Asia* → in scope. Where it builds does not change what it is. - *Acquired by a larger group but still operating as a distinct product line* → out of the ranking, noted in the annex. Its funding trajectory stopped being a signal on acquisition.

Why the boundary matters: an interviewer will test it with an edge case. This definition is built to survive "what about company X?" — every exclusion above has a stated reason, and the reason is always the same one: the criteria must measure comparable things.

2. Company list

26 companies. Funding figures are deliberately absent here — collecting and sourcing them is the job of steps 1 and 2. This list is names, geography and segment only.

Sources used to build it: the Swiss Robotics Startups ecosystem map (~100 entries, filtered against the section 1 criteria), Greater Zurich Area's robotics selection, Top100 Startups Switzerland, and EU-Startups / Tech.eu funding coverage from 2025–2026.

#	Company	Country	Sub-segment	How I found it
1	ANYbotics	CH — Zurich	Legged robots, industrial inspection	Known ETH spin-off
2	Verity	CH — Zurich	Autonomous warehouse inventory drones	Ecosystem map
3	Voliro	CH — Zurich	Contact-inspection drones	Ecosystem map
4	Flyability	CH — Lausanne	Confined-space inspection drones	Known EPFL spin-off
5	Wingtra	CH — Zurich	VTOL mapping and survey drones	Ecosystem map
6	mimic robotics	CH — Zurich	Dexterous manipulation	EU-Startups, Nov 2025
7	Ascento	CH — Zurich	Wheeled-legged security patrol robots	Ecosystem map
8	Gravis Robotics	CH — Zurich	Autonomous heavy construction machinery	Greater Zurich selection
9	LOXO	CH — Bern	Autonomous delivery vehicles	Ecosystem map
10	Ecorobotix	CH — Yverdon	Precision agriculture, smart spraying	Ecosystem map
11	Tethys Robotics	CH — Zurich	Autonomous underwater robots	Greater Zurich selection
12	Saeki	CH — Zurich	Robotic fabrication for construction	Ecosystem map
13	NEURA Robotics	DE — Metzingen	Cognitive robots and humanoids	Largest EU round, 2026
14	Agile Robots	DE — Munich	Force-controlled manipulation	Sector coverage
15	Sereact	DE — Stuttgart	Embodied AI for robotic picking	EU funding coverage 2026
16	Robco	DE — Munich	Modular industrial robots for SMEs	Sector coverage
17	Fernride	DE — Munich	Autonomous yard and port logistics	Sector coverage
18	Humanoid	UK — London	Industrial humanoid robots	Tech.eu, Jul 2026
19	Dexory	UK — Oxford	Warehouse inventory robots	Sector coverage
20	Wayve	UK — London	Embodied AI for driving — see edge case	Sector coverage
21	Exotec	FR — Lille	Warehouse goods-to-person robotics	Sector coverage
22	Donecle	FR — Toulouse	Automated aircraft inspection drones	EU-Startups, Apr 2026
23	AgreenCulture	FR — Toulouse	Autonomous agricultural robots	Sector coverage
24	PAL Robotics	ES — Barcelona	Humanoid and service robots	Sector coverage
25	THEKER	ES — Barcelona	Humanoid manipulation	EU funding coverage 2026
26	Oversonic Robotics	IT — Besana Brianza	Humanoid robots for industry	Sector coverage

2b. Edge cases — ruled on before scoring

Applying section 1 to the grey zone. Each ruling is stated so it can be defended.

Company	Ruling	Reason
Auterion (CH)	Out	Drone autonomy software running on third-party airframes. No own physical product.
Energy Robotics (DE)	Out	Inspection software layered on other vendors' robots. Same reason as Auterion.
Bota Systems (CH)	Out	Force-torque sensor supplier — a component, not a system that perceives and acts.
Distalmotion (CH)	Out	Surgical robotics. Strong Swiss player, but medtech regulation and sales cycles make it incomparable to industrial Physical AI.
Wayve (UK)	In, flagged	Does not build the vehicle, which strains the "physical product" test. Kept because the autonomy stack <i>is</i> the product and it is deployed on real roads. Reversible if the flag proves distorting.

The Wayve ruling is the one to expect a question on. The honest answer is that it sits on the boundary, the boundary was drawn before the company was considered, and the decision is documented rather than silent.

3. Scoring criteria and weights

Weights must sum to 100. The justification column is what gets discussed in an interview — it is more important than the number itself.

Criterion	Weight	What it measures	Why this weight
Field traction	30	Units deployed at paying industrial customers. Funding is a secondary indicator only.	In Physical AI the binding constraint is deployment, not modelling. Anyone can demo; almost nobody survives contact with a real site.
Team / execution	25	Founder background, hardware iteration speed, ability to hire scarce robotics talent	Hardware punishes slow teams. Iteration speed is the single best predictor of who is still standing in three years.
Technology / product	20	Maturity, technical moat, defensibility	Weighted below traction on purpose: robotics moats erode fast because everyone reads the same papers. What defends is accumulated real-world data, which shows up under traction.
Market	15	Addressable size, segment growth rate	Discriminates weakly here — most of the panel targets comparable industrial segments, so it separates few companies.
Timing	10	Why now — what changed that makes this possible today	Deliberately low. See section 3b.
Total	100		

3b. Two positions this weighting takes

Traction is not the amount raised. It is redefined as units deployed at paying industrial customers, with funding demoted to a secondary indicator. Dollars raised is a lagging, noisy signal: a €100M round reports what a group of investors believed eighteen months ago. Robots running on a customer site report what works now.

Timing is capped at 10 on purpose. The "why now" of Physical AI — cheap sensors, vision- language-action models, industrial labour shortage — is identical for every company in the panel. A criterion that awards everyone the same score does not rank anything; it only inflates the totals. Keeping it at 10 acknowledges that it matters to the sector without pretending it separates companies within it.

4. Scoring scale

A 1–5 scale with written anchors for every criterion, so that two runs of the pipeline produce comparable scores. Without anchors the model drifts between runs and the ranking means nothing.

Scores 2 and 4 are intermediate positions between the anchors below.

Field traction — weight 30

Score	Anchor
1	Demos, pilots and letters of intent only. No customer paying for a deployment, or deployments only at investor and partner sites.
3	Paying deployments at roughly 3–10 customers, still sold project by project, each installation substantially custom.
5	Repeat orders from several industrial customers, tens of units in routine daily service, and at least one customer that has moved past an initial pilot into a fleet.

*The discriminator is the **pilot-to-fleet transition**, not the number of pilots. Robotics is full of companies with fifty pilots and no second order — that is "pilot purgatory", and it scores a 2, not a 4. A single customer that scaled from three units to thirty is worth more evidence than twenty logos on a website.*

Team / execution — weight 25

Score	Anchor
1	First-time founders with no record of shipping hardware, no visible iteration, flat or shrinking team.
3	Credible technical founders with a strong research pedigree, one product generation shipped, steady hiring.
5	Founders who have shipped hardware at scale before, or two-plus product generations shipped at this company with visibly shortening cycles, and the pull to hire senior robotics people away from big tech.

*An **ETH or EPFL pedigree caps at 3 in this panel**. Nearly every Swiss company on the list has one, so it is table stakes, not a differentiator. Anything that everyone shares cannot separate anyone — the same logic that caps the timing criterion at 10.*

Technology / product — weight 20

Score	Anchor
1	The capability is demonstrated in published research by others and the demo could be reproduced with off-the-shelf components.
3	Real engineering depth — hard integration work, reliability in the field — but the underlying approach is broadly available to competitors.
5	An advantage that compounds: a proprietary real-world data flywheel, a hardware component competitors cannot buy, or a certification that takes years to replicate.

*A 5 requires something that **gets stronger with use**. A clever architecture does not qualify; everyone reads the same papers. This is the criterion where the weighting argument from section 3b cashes out.*

Market — weight 15

Score	Anchor
1	Narrow niche with few buyers, or a segment where no budget line exists for this category of spend.
3	Genuinely sizeable segment, but crowded, or with procurement cycles long enough to starve a startup.
5	Large segment with an urgent and already-budgeted pain — labour shortage, regulatory inspection mandate — where the buyer purchases in this category today.

*Judge the **existing budget line**, not the headline market size. A TAM is easy to inflate and nobody can check it. A buyer who already has a line item for inspection is a buyer who can sign.*

Timing — weight 10

Score	Anchor
1	Could have been built five years ago and was not, with no external change explaining why now.
3	Rides the general Physical AI wave — the same "why now" as every other company in the panel. This is the default score here.
5	A specific, dateable unlock the company is positioned for: a regulation entering into force, a component crossing a price threshold, a customer-side mandate.

Because the generic wave scores a 3 for everyone, only a specific unlock earns a 5. This is what keeps a low-weight criterion from quietly inflating every total.

4b. Handling missing evidence

The pipeline will not find complete data for all 26 companies. Younger and less publicised companies will have thinner public footprints, and that must not be silently rewarded or punished by guesswork.

Rule: where evidence for a criterion is absent, the score is capped at 2 and flagged as low confidence. The model must never infer a score from what is plausible for a company of that type.

Absence of evidence is recorded as absence, not averaged away into a middling score. Every low-confidence flag is listed in the final report so a reader can see exactly which parts of the ranking rest on thin data.

4c. Sourcing the traction criterion

There is a real tension in this design, and it is better stated than discovered: field traction carries the highest weight (30) and is the hardest evidence to obtain publicly. Industrial buyers rarely disclose deployment volumes. Left alone, the most important criterion would carry the most low-confidence flags.

The weight is not lowered — the argument for it holds. Instead, step 1 searches for traction evidence deliberately and beyond the company's own marketing:

Source	What it reveals
Customer-side press releases	The buyer announcing a rollout is stronger evidence than the vendor announcing a win
Published case studies	Usually name a site and sometimes a unit count
Job postings	A company hiring ten field deployment engineers has deployments. Hiring only researchers does not.
Trade and industry press	Covers installations that never reach the startup press
Regulatory and tender records	Public-sector deployments leave a paper trail

Job postings are an **indirect indicator** and are labelled as such in the extracted data. They support a score, they never establish one on their own.

4d. Model allocation across the pipeline

Two models, assigned per step rather than one model for the whole pipeline. The principle: **spend capability where judgment happens, not uniformly.**

Step	Model	Reasoning
1. Source collection	claude-sonnet-5	High input volume (search results), low judgment. Finding and returning pages is not where model capability separates.
2. Structured extraction	claude-sonnet-5	Reading a page and filling a schema field with its source URL. Mechanical once the schema is fixed.
3. Weighted scoring	claude-opus-5	The only step that is pure judgment: arbitrating between the 1/3/5 anchors, applying the missing-evidence cap, and writing a justification that survives challenge. This is the step the project is actually about.
4. Report generation	none	Formatting already-decided content — and once that is true, a model is the wrong tool. Assembled in Python from the stored records, so no figure can drift from the record it came from. Free and reproducible. Sonnet was allocated here originally; the change is recorded rather than silent.
5. Thesis and executive summary	claude-opus-5	Synthesis across 26 scored companies — pattern-finding, not transcription.

Cost at August 2026 list prices: Opus is 2.5× Sonnet per token on both input and output. Running the whole pipeline on Opus would roughly double the project cost for a quality gain confined to steps 3 and 5.

Consequence to respect on any rerun: step 3 stays on Opus. A scoring pass run on a different model is not comparable to the existing one, and the calibration in §5 would have to be redone.

5. Calibration set

Five companies scored **by hand, before running the pipeline**. The gap between these scores and the model's output is measured and discussed in the final report.

This is the section that makes the project credible. Do not skip it.

Hand-scoring completed 2026-08-18, blind, before any pipeline code was written. Full per-criterion scores, confidence levels, and justifications are in `calibration-worksheet.md`.

Scored 2026-08-18 on `claude-opus-5`. The model saw only the extracted evidence: never the open web, never these hand-scores, and never another company. It did not compute its own total — the weighting is applied in `src/scoring.py`, by the same arithmetic as the worksheet.

Company	My score	Model score	Gap	Who was right, and why
ANYbotics	93/100	86/100	-7	Model, narrowly, on technology. I scored the moat 5; it argued 4 by conceding the ATEX/IECEx Zone 1 certification and then pricing the IP against it — 8 patents filed, a 2019 patent licence paid to Boston Dynamics, and the ANYdrive actuator now built by an outside supplier. That is a specific case I did not make. Traction, team and timing agree exactly.
Verity	92/100	85/100	-7	Draw, same shape as ANYbotics. Identical pattern: agreement on traction and team, one notch down on technology and market. Two companies I rated within a point of each other are still within a point of each other.
Gravis Robotics	85/100	70/100	-15	Model, on technology (-2). I scored 5. It found no compounding advantage: commodity cameras, lidar and GNSS, no certification barrier, BuiltWorlds placing Gravis beside Hive Autonomy and SafeAI doing the same retrofit approach, and the CEO's own point that electronic joystick signals make retrofits easy — which opens the same door to competitors. I scored this the day after a \$200M round; I should treat that as a possible influence on my read.
Humanoid	77/100	59/100	-18	Model, and the test inverted — see below. Traction -2 and market -2 carry the whole gap.
mimic robotics	60/100	54/100	-6	Draw, and the closest agreement in the set. Traction identical at 2. It is +1 on market where I was harsh, -1 on technology and team. Nothing here separates us.

Agreement across the 25 criterion scores: **12 exact, 22 within one notch**, three disagreements of two notches (Gravis technology, Humanoid traction, Humanoid market).

The model is stricter than me on every company, never more generous. A directional bias is a more useful finding than scatter would have been: it is one correction to argue about, not five.

The ranking is identical. ANYbotics, Verity, Gravis Robotics, Humanoid, mimic robotics, in that order, on both sides. Where the totals move, the order does not.

Where the disagreement actually sits, summed across the five companies in notches:

Criterion	Sum of gaps	Reading
Team / execution	-1	Near-consensus. We read founders and hiring the same way.
Timing	+1	Consensus, helped by the 3 default doing its job.
Field traction	-3	Concentrated: Humanoid -2 and Gravis -1. The other three agree exactly.
Market	-4	The model discriminates where I did not. I gave four 5s (\$5, third observation, called this out as a flaw in my own scoring). It gave 4/4/4/3/3.
Technology	-5	The real disagreement. It reads moats as thinner than I do, on four of five companies.

Technology is the argument to prepare. My scores lean on engineering depth; its scores ask what *compounds* — which is closer to what the anchor actually says. That is worth conceding or defending explicitly, not splitting.

Three observations from the hand-scoring pass, recorded before seeing any model output:

- **The Humanoid test no longer isolates what it was designed to test.** It was chosen to check whether the model conflates funding with traction — traction was expected to score low. It scored 4, on the strength of a dated, binding Schaeffler contract. The reasoning holds, but if the model also scores it high, the two possible causes (reading the contract vs. being swayed by headline numbers) become indistinguishable.

Resolved by the scored run, in the opposite direction to the one the test anticipated. The model scored Humanoid's traction **2**, against my 4 — it was the stricter reader, not the looser one. It credits the Schaeffler agreement for lifting the score above a pure 1, then refuses 3 because there are no paying installs at three to ten customers, and it flags the "nine industrial deployments with Fortune 500 customers" as single-sourced, self-reported and internally inconsistent, and the 34,000 pre-orders as unverified. On the evidence in front of it, that reading of the anchor is better than mine: I let a binding future commitment stand in for present deployment. This one disagreement is -12 of the -18 point gap on Humanoid. - **Gravis Robotics no longer tests the coverage floor.** It was chosen for near-zero public coverage. A \$200M SoftBank round announced 2026-08-17 changed that overnight; it scored 85 with HIGH confidence on four of five criteria. **mimic robotics is now the only functioning missing-evidence test** — it produced the single LOW-confidence flag in the set. - **Market does not discriminate.** Four of five companies scored 5. That is the same failure the weighting section identifies for Timing, and it suggests the Market "5" anchor is too easy to reach. Weight is only 15, so the effect on the ranking is bounded — but it is a real finding about the criteria, not about the companies.

5b. Observations from the extraction pass, before any scoring

Step 2 ran on all five companies on 2026-08-18. Two findings are recorded here, before the scores exist, because they change how the table above should be read.

- **The missing-evidence test did not fire on the company it was built for.** mimic robotics is the only company in the set still carrying a thin public footprint, and it was hand-scored 60/100 largely for that reason. Extraction returned `field_traction: found` at HIGH confidence. Its step 1 report states

plainly that no named paying customer was confirmed by any third party. Neither `searched_not_found` nor `not_searched` was assigned, so the §4b cap never engaged.

- **A single item carried that rating, and it is not traction evidence.** HIGH requires one direct claim from an independent source. For mimic robotics that claim is a July 2026 trade article reporting a new video-action model developed with Black Forest Labs — a product announcement naming no customer, site, unit count, or agreement. The other six items are all `indirect`. The same count is 14 for Verity, 9 for ANYbotics, 8 for Gravis Robotics and 4 for Humanoid. All five are rated HIGH.

Corrected after scoring — the warning above applies to step 2 only. Written before step 3 ran, this section concluded that the pipeline had not measured mimic robotics' weakness. The scored result contradicts that. Opus rated its field traction **2/5, the same value I gave it by hand**, and its justification names precisely what the confidence flag missed: *every customer claim is unnamed and company-sourced*, SiliconANGLE *explicitly notes no names were disclosed*, and nothing shows *a paying deployment, a unit count, or a repeat order*.

So the gap is real but it is contained. Step 2's confidence rule was fooled; step 3 was not, because it reads `evidence_grade` and `attributed_to` rather than the summary flag. The honest statement is narrower than the one first recorded here: **a HIGH confidence flag is not a statement about how strong the evidence is, and should not be read as one** — in the report, or by anyone using this pipeline. The scores themselves are not affected.

6. Known limitations

Stated upfront, not discovered by the interviewer.

A search budget manufactures absence. The first ANYbotics collection ran with a budget of 12 searches and concluded that no customer-side press release, tender record, or regulatory filing named the company. Re-run at 20 with the same prompt, it found four — Outokumpu, PETRONAS, Equinor, GE Vernova — plus patent data it had previously reported as unsearched. Under §4b that difference is not cosmetic: absent evidence caps a score at 2, so **the search budget was an input to the ranking, not only to the cost.** Two mitigations are in place: the budget is now 20, and step 2 records `searched_not_found` separately from `not_searched` so that only genuine absence triggers the cap. Neither proves sufficiency — 20 searches are still fully consumed on data-rich companies, so the ceiling may still bind on the largest ones. Every capped score should be read as *not found within 20 searches*, never as *does not exist*.

Confidence measures the kind of source, not the weight of the evidence. It is computed in Python, not asked of the model: HIGH requires at least one direct claim from an independent source (customer-side, tender record, or trade press), MEDIUM is everything else with evidence, LOW is absence. This makes §4b enforceable and removes run-to-run drift, at a stated cost — one independent source and six both return HIGH. Confidence here answers *who said it*, not *how much of it there is*.

Narrowed after the first run. HIGH now requires two independent direct sources rather than one, so it means *corroborated*. Seven of the flags across the five companies moved down as a result. The limitation is real but smaller than first written: above two, volume still does not move the flag.

Nothing checked that a piece of evidence belonged to the criterion it was filed under.

Extraction assigns each item to a criterion, and the confidence rule then counts items; neither step asks whether the claim is *about* that criterion. Measured on mimic robotics (§5b): a product announcement filed under field traction was graded a direct claim from an independent source and, on its own, lifted the company's weakest criterion to HIGH. Combined with the binary rule above — one qualifying item scores the same as fourteen — a single misfiled item is enough to erase a coverage gap. Reading the evidence list is part of reading the score here, not an optional audit on top of it.

Addressed, and the fix is not free of its own assumption. An evidence item must now name what makes it direct — a named customer, a unit count, a signed agreement, a dated deployment, or a regulatory record — and one that names none of those is recorded as indirect regardless of the grade the model assigned. The scoring prompt also now states that one fact scores under one criterion only, with the ATEX case named. Both rules take effect on the next run of their step; the scores in §5 predate them. The assumption they rest on is that the model labels its own reasoning honestly, which is weaker than a check that reads the claim itself.

Public evidence only. No interviews, no customer references, no proprietary databases. This favours companies that publicise well over companies that deploy well — exactly the confusion §3b sets out to avoid. Searching customer-side sources rather than vendor announcements (§4c) reduces the bias; it does

not remove it. The step 1 prompt also does not explicitly request German-, French-, or Italian-language sources, which matters for a panel centred on Switzerland. Whether that changes coverage materially has not been measured.

Errors in step 1 are invisible to step 3. The scoring model never sees the web — it scores the evidence document that collection produced. A source that was missed, or misread, propagates silently into the ranking with no mechanism downstream to catch it. The audit trail exists (every claim carries its URL) but auditing it is a manual act, not an automated check.

The ranking is a snapshot, and the panel moves faster than the snapshot. Gravis Robotics went from near-zero public coverage to a \$200M SoftBank round announced 2026-08-17 — during the scoping work itself. Any score in this report is a statement about the evidence available on its collection date, which is recorded per company.

The calibration set is five companies scored by one person. It measures agreement between a hand-scorer and the model, not correctness of either. Two of the five have also lost the property they were chosen to test: Humanoid no longer isolates funding-versus-traction, and Gravis no longer tests the coverage floor (§5). mimic robotics is the only functioning missing-evidence test left, so the missing-evidence machinery rests on a sample of one.

A quarter of the weight sits on criteria that barely separate the panel. Timing is capped at 10 because the "why now" is identical sector-wide (§3b) — that is deliberate. Market at 15 was expected to discriminate and, on the calibration set, did not: four of five companies scored 5. Together that is 25 points of 100 doing little ranking work, which concentrates the real discrimination in traction and team. Raising the Market bar is a change to consider after the full run, not mid-flight.

One fact scores twice. ANYbotics' ATEX/IECEEx Zone 1 certification is credited under technology, as the multi-year certification moat the 5 anchor names, and again under timing, as a dated external unlock. It is the single most important fact in the set, and the rubric counts it in two places, which inflates the company that holds it. Found by the pipeline's own synthesis step rather than by me, which is the reason it is written here instead of quietly patched: a criterion set that can double-count is a finding about the criteria. The fix is to decide which criterion owns a certification before the next run, not to adjust one score after the fact.

7. Annex — ecosystem players, not scored

Excluded from the ranking because the scoring criteria do not apply to them, but they shape the sector and belong in the picture.

Foreign corporate R&D sites in Switzerland

Not scored — the criteria are built for independently financed companies. Listed because they compete for the same talent pool as every company in the ranking, and because they are where a share of Swiss robotics and computer-vision work actually happens.

Group	Site	Relevance to Physical AI
Google / DeepMind	Zurich	Largest Google R&D centre outside the US; robotics and computer vision
Amazon	Zurich	Reinforced by the RIVR acquisition, see below
Apple	Zurich	AI and computer vision lab
Meta	Zurich	Computer vision
Microsoft	Zurich	Mixed reality and AI
NVIDIA	Zurich	Robotics and simulation
Disney Research	Zurich	Character robotics, long-standing ETH partnership
IBM Research	Rüschlikon	Long-standing European research site
ABB	Zurich region	Not foreign, but a listed incumbent — excluded from the ranking for scale

To verify and complete in step 1 — headcount and current robotics activity per site.

Acquired companies

The most interesting rows in this document. Both were ETH-originated, both were acquired before reaching independent scale.

Company	Acquirer	Year	Signal
RIVR (ex Swiss-Mile)	Amazon (US)	March 2026	Wheeled quadrupeds for last-mile delivery. Raised ~\$25M incl. Bezos Expeditions before exit.
Sevensense Robotics	ABB (CH/SE)	2024	Visual navigation for mobile robots. Absorbed into an incumbent's product line.

European companies that relocated their HQ abroad

Company	Founded in	Moved to	Year
---------	------------	----------	------

To investigate in step 1

Working hypothesis for the thesis

Emerging from the annex, **to be tested against the scored data, not assumed**: Switzerland generates world-class Physical AI technology through ETH and EPFL spin-offs, but value capture migrates abroad — through acquisition by foreign strategics — before those companies reach independent scale. If the scored data supports it, the interesting question becomes which of the 26 has the field traction to break that pattern.